

Install and activate a virtual environment running R in Jupyter notebook.

Run the R tutorial notebook `Beginning R.ipynb` and study its contents. A further resource is available in the Supplemental Materials section.

Using the Kaggle dataset (you may need to sign up for Kaggle if you have not already done so), download the data for the house prices competition.

<https://www.kaggle.com/c/house-prices-advanced-regression-techniques/data>

Keep the following columns:

"Id", "LotFrontage", "LotArea", "Neighborhood", "OverallQual", "RoofStyle", "YearBuilt", "YearRemodAdd", "MoSold", "YrSold", "SaleType", "SaleCondition", "SalePrice", "BedroomAbvGr"

Show a summary of your dataset using `str`.

Get the average house price.

Using the `plot()` command, plot all variables against each other. Take note of any interesting correlations and do further plots on those pairs.

Plot the dimensions OverallQual, SalePrice, Neighborhood and BedroomAbvGR on the same chart.

Using the `boxplot` function or `barplot` function, show the distribution over the variables LotFrontage, LotArea, Neighborhood, OverallQual, RoofStyle, YearBuilt, YearRemodAdd, YrSold, SaleType, SaleCondition, SalePrice, BedroomAbvGr showing them all in a 3x3 square of charts.